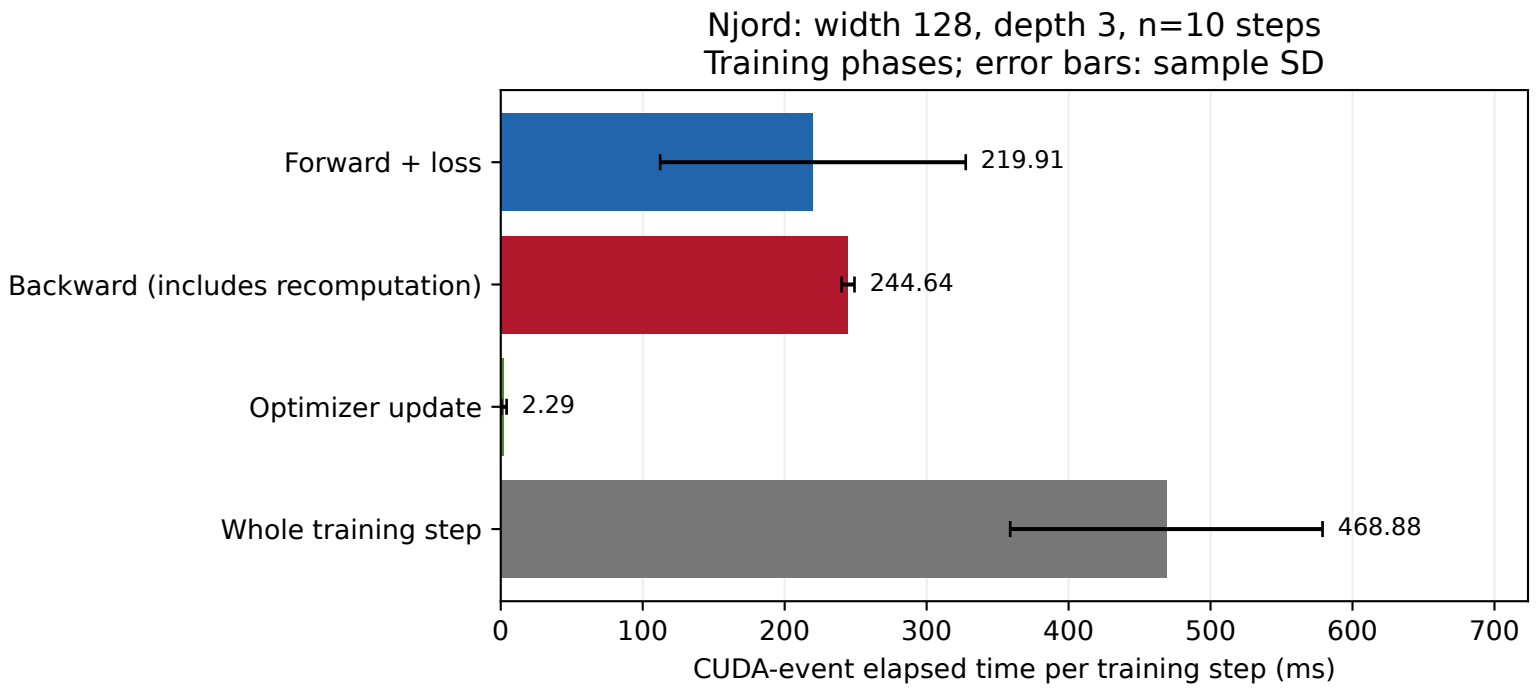
Njord: width 128, depth 3, n=10 steps
Training phases; error bars: sample SD
Forward + loss
219.91
Backward (includes recomputation)
244.64
Optimizer update
2.29
Whole training step
468.88
0
100
200
300
400
500
600
700
CUDA-event elapsed time per training step (ms)