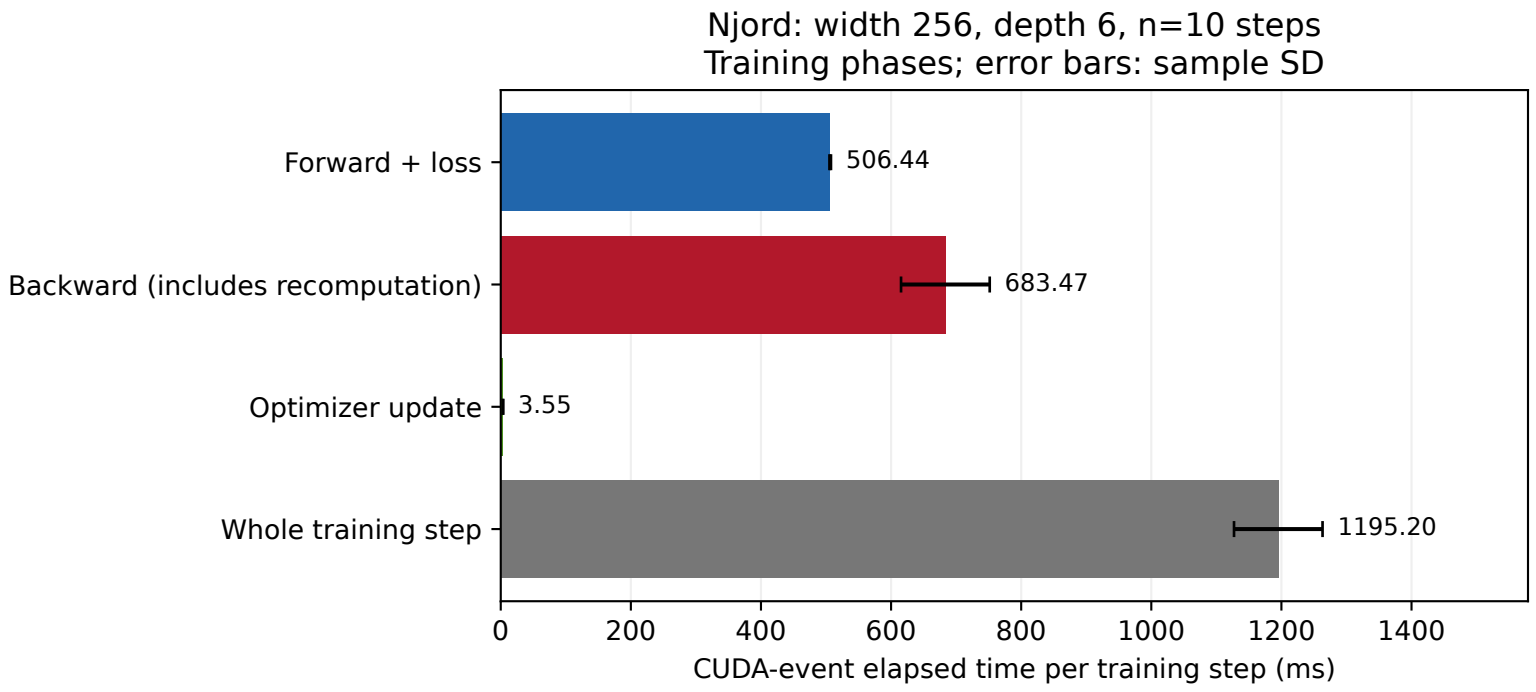
Njord: width 256, depth 6, n=10 steps
Training phases; error bars: sample SD
Forward + loss
506.44
Backward (includes recomputation)
683.47
Optimizer update
3.55
Whole training step
1195.20
0
200
400
600
800
1000
1200
1400
CUDA-event elapsed time per training step (ms)